

Agentic Task Force Canvas – Full Feature List

1. Workflow Design & Orchestration

- Visual drag-and-drop canvas to design multi-agent workflows
- Sequential, parallel, conditional, and looping execution flow
- Color-coded routing lines indicating execution condition (success, failure, retry, keyword/context-based)
- Connectors with built-in routing logic configuration (dropdowns)
- Agent linking logic for fallback, retries, or data-dependent branching
- Global preview of execution paths and failure handling

2. Agent Configuration & Management

- Configurable agent blocks with editable name, description, and context role
- Agent tool selection (API access, DB queries, document tools, RAG modules)
- Memory configuration (ephemeral, persistent, vector-based)
- Reusable agent templates stored and accessible in a left-hand library
- LLM model selector per agent (OpenAI, Claude, Mistral, local LLMs)
- Graphical interface for tool/data/memory linking per agent

3. Knowledge Integration & Querying

- Procurement Knowledge Graph connectivity (Neo4j or AWS Neptune)
- Visual linking of agents to KG queries and relationships
- Embedded Cypher and GraphQL template support
- Keyword and context-condition matching using knowledge graph values
- Entity navigation: Supplier → Contracts → SLA → Invoices
- Real-time query preview and semantic query augmentation

4. Routing Intelligence & Logic

- Routing Inspector with dropdown condition builder (success/failure/output

contains/context match)

- Context-aware routing with category assignment (e.g., invoice, dispute, contract)
- Confidence threshold support for routing decisions (e.g., 85% match → route)
- Learning-enabled routes that adapt based on historical outcomes
- Human-correctable routing paths via natural language override
- Conditional pathing and retry configuration per connector

5. Execution Monitoring & Logging

- Live execution view with step-by-step agent status display
- Retry, error, and fallback indicators
- Structured execution logs per run with route taken, input/output trace
- Visual highlights of active and completed agent flows
- Agent-specific run metrics and summaries

6. Learning & Feedback Loop

- Feedback capture from human reviewers (text input for correction)
- Route learning: agent paths improved through feedback classification
- Auto-updated routing suggestion engine based on interaction logs
- Score adjustment and training dataset generation for fine-tuning LLM routing
- Visual audit trail of feedback and outcomes

7. Security, Tenancy, and Governance

- Multi-tenant deployment model with per-tenant namespace isolation
- Role-based access control (RBAC) per workflow and agent type
- Data masking and policy tagging for sensitive data
- Per-agent permissioning for data access and actions
- Encrypted audit trail logging with searchable event history